

# Real-Time Fraud Detection in Transactions.

---

## 1. DATASET: THE PAYSIM DATASET

### 1.1 Background and Origin

The **PaySim dataset** is a **synthetic dataset** that simulates **mobile money transactions**. Originally, it was published on **Kaggle**, where it can be downloaded either manually or via the **Kaggle API**. The dataset is designed to **mimic real financial transactions** in terms of scale, distribution, and types of operations, while not revealing any sensitive private customer data (since it's artificially generated).

#### Why Synthetic?

Real financial transaction data is typically **strictly confidential** (due to privacy regulations and to avoid revealing vulnerabilities). Hence, PaySim was created to provide **researchers** and **developers** a near-realistic scenario to test **fraud detection** algorithms.

#### Kaggle Link:

We can find the dataset on Kaggle by searching for "PaySim dataset." Typically we can do something like:

```
kaggle datasets download -d YOURUSERNAME/paysim
```

to retrieve it with the Kaggle CLI, or we can download it directly from the Kaggle website.

### 1.2 Size and Format

- **Rows:** ~6.36 million transactions (some versions might slightly differ).
- **File Size:** ~500 MB in CSV form.

- **Format:** Plain CSV with a header row.

This large volume makes it ideal for **big data** or **streaming** pipeline demonstrations.

## 1.3 Columns in the PaySim Dataset

The dataset typically includes the following columns (naming might differ slightly depending on the version used):

### 1. **step**

- *Integer:* Each "step" corresponds to a discrete time unit in which transactions occurred.
- Think of it as an approximation of "hours" or "minutes," but the dataset's authors just named it "step."

### 2. **type**

- *Categorical:* Transaction type. Common values include:
  - **PAYMENT**
  - **TRANSFER**
  - **CASH\_OUT**
  - **DEBIT**
  - **CASH\_IN** (in some versions)
- This is an important feature because **fraud** often correlates with **TRANSFER** or **CASH\_OUT**.

### 3. **amount**

- *Float:* The monetary value of the transaction. Ranges from small amounts to very large ones (up to  $10^7$  or so).

### 4. **nameOrig**

- *String:* The unique identifier of the sender (origin) account. Often in the format **Cxxxxxxx**.

### 5. **oldbalanceOrg**

- *Float*: The sender's account balance before the transaction.

## 6. newbalanceOrig

- *Float*: The sender's account balance after the transaction.

## 7. nameDest

- *String*: The unique identifier of the recipient (destination) account. Often Cxxxxxxx for customer, or Mxxxxxxx for merchant.

## 8. oldbalanceDest

- *Float*: The recipient's account balance before the transaction.

## 9. newbalanceDest

- *Float*: The recipient's account balance after the transaction.

## 10. isFraud

- *Integer (0 or 1)*: **Primary label** indicating whether the transaction is fraudulent (1) or legitimate (0).
- Highly imbalanced: only a tiny fraction are 1.

## 1. isFlaggedFraud

- *Integer (0 or 1)*: Whether the transaction was flagged as fraud by some older heuristic approach.
- Typically near zero in most rows.

Dataset loaded. Shape: (6362620, 11)

	step	type	amount	nameOrig	oldbalanceOrg	newbalanceOrig	nameDest	oldbalanceDest	newbalanceDest	isFraud	isFlaggedFraud
0	1	PAYMENT	9839.64	C1231006815	170136.00	160296.36	M1979787155	0.0	0.00	0	0
1	1	PAYMENT	1864.28	C1666544295	21249.00	19384.72	M2044282225	0.0	0.00	0	0
2	1	TRANSFER	181.00	C1305486145	181.00	0.00	C553264065	0.0	0.00	1	0
3	1	CASH_OUT	181.00	C840083671	181.00	0.00	C38997010	21182.0	0.00	1	0
4	1	PAYMENT	11668.14	C2048537720	41554.00	29885.86	M1230701703	0.0	0.00	0	0
5	1	PAYMENT	7817.71	C90045638	53860.00	46042.29	M573487274	0.0	0.00	0	0
6	1	PAYMENT	7107.77	C154988899	183195.00	176087.23	M408069119	0.0	0.00	0	0
7	1	PAYMENT	7861.64	C1912850431	176087.23	168225.59	M633326333	0.0	0.00	0	0
8	1	PAYMENT	4024.36	C1265012928	2671.00	0.00	M1176932104	0.0	0.00	0	0
9	1	DEBIT	5337.77	C712410124	41720.00	36382.23	C195600860	41898.0	40348.79	0	0

# 1.4 Key Dataset Characteristics

- **Synthetic** but distributionally realistic.
  - **Unbalanced**: ~0.13% or fewer transactions labeled as fraud.
  - **Transaction diversity**: multiple transaction "types," with typical money-laundering or fraud patterns lying mostly in **TRANSFER** and **CASH\_OUT**.
- 

## 2. EXPLANATION OF EACH COLUMN

Although summarized above, let's detail each column thoroughly:

### 1. **step**

- Each unit in this feature could be considered an incremental time step. Some interpreters assume it's an hour. If the highest step is ~743, that might represent a month's worth of hours.
- For time-series or streaming analysis, **step** can be used to see how transactions evolve over time.

### 2. **type**

- Common categories:
  - **TRANSFER**: Money moves from one account to another.
  - **CASH\_OUT**: The account holder withdraws money, presumably as cash.
  - **PAYMENT**: Payment to a merchant.
  - **DEBIT**: Possibly a direct debit from an account by a third party.
  - (Sometimes **CASH\_IN** if the dataset variant includes it.)
- This is an important predictor because historically, fraudulent transactions often cluster in certain types.

### 3. **amount**

- The transaction's monetary value. Ranges widely from negligible amounts to multi-million.
- Fraud can be high or low amounts, so we must carefully consider distribution.

#### 4. nameOrig / nameDest

- Strings: Cxxxxxxx for customer, Mxxxxxxx for merchant.
- Not always needed for direct numeric modeling but can be used for feature engineering (like counting repeated patterns to the same "destination," etc.).

#### 5. oldbalanceOrg / newbalanceOrg

- The balance before and after the transaction for the origin's account.
- Potentially important: a fraudulent user might have suspicious patterns like "oldbalanceOrg = amount, newbalanceOrg = 0" repeatedly.

#### 6. oldbalanceDest / newbalanceDest

- Similarly for the recipient's account.
- If a destination is frequently receiving large amounts, that might be a red flag or normal merchant activity.

#### 7. isFraud

- The main label used for **supervised** or **evaluation** tasks.
- In our anomaly detection approach, we often used this label **only** at the evaluation stage, because we do an **unsupervised** or **semi-supervised** approach.

#### 8. isFlaggedFraud

- Often zero in most rows.
- The original creators said it flags transactions that are abnormally high (some threshold). Not extremely useful in practice, but can be a baseline.

---

## 3. OBJECTIVE OF THE PROJECT

**Goal: Real-Time Anomaly Detection in Financial Transactions** using the PaySim dataset as a stand-in for real mobile money data. Specifically:

1. **Detect** suspicious or fraudulent transactions **on the fly** (near real time).

2. **Leverage** an **unsupervised** or **semi-supervised** approach (e.g., IsolationForest, LocalOutlierFactor, Autoencoders) to handle the extreme class imbalance.
  3. **Demonstrate** a big data streaming pipeline with **Kafka + Spark Structured Streaming** for real-time ingestion.
  4. **Provide** a user-friendly dashboard with **Streamlit** that can visualize the flagged anomalies in near-real time, offering **insights** to risk analysts or external stakeholders.
- 

## 4. INTRODUCTION TO THE PROJECT

### 4.1 Overview

This project aims to simulate a **full end-to-end** pipeline for **fraud detection**:

1. **Data Ingestion**: Instead of a real transaction feed, we have the PaySim CSV. We mimic streaming by publishing rows to **Kafka**.
2. **Stream Processing**: **Spark Structured Streaming** consumes from Kafka, runs a **pre-trained** anomaly detection model (like IsolationForest), and flags suspicious transactions.
3. **Storage**: The flagged anomalies are written out to **CSV** files in near-real time.
4. **Visualization**: A **Streamlit** dashboard loads those flagged anomalies, draws charts, and provides a user-friendly monitoring tool.

### 4.2 Contents of the Project

1. **EDA & Feature Engineering Notebook** (a local or Colab notebook)
2. **Model Training** with unsupervised ML (IsolationForest, LOF, Autoencoder).
3. **Kafka Producer Script** ( `kafka_producer.py` ): Streams rows from the CSV into Kafka topic.
4. **Spark Streaming Script** ( `spark_streaming.py` ): Subscribes to Kafka, loads pre-trained model, flags anomalies, writes them to CSV.

5. **Streamlit Dashboard** ( `dashboard.py` ): Displays flagged anomalies in real time.
6. **Deployment**: GCP VM (Compute Engine) used to host Kafka + Spark + the entire pipeline.

## 4.3 Why PaySim?

- The PaySim dataset is large (~6 million rows), capturing the complexity of real transactions.
- The presence of an **isFraud** label allows for **evaluation** of anomaly detection.
- It's free, synthetic, and widely known in academic/industry examples.

## 4.4 Why Real-Time?

Fraud detection is most valuable when **immediate**. The faster we detect suspicious activity, the faster we can block it or investigate. Traditional offline batch detection might be too slow.

---

# 5. EXPLORATORY DATA ANALYSIS (EDA), FEATURE ENGINEERING, AND MODEL TRAINING

## 5.1 EDA

### 5.1.1 Basic Statistics

- **Total rows**: ~6,362,620
- **Fraud**: ~8,213 (only ~0.13%).
- **Transaction Types**: Payment, Transfer, Cash Out, Debit, etc.
- **Distribution of `amount`**: Skewed; many small transactions, few extremely large.

### 5.1.2 Visualizations

#### 1. Fraud vs. Non-Fraud Count


- A bar chart showing an enormous imbalance: ~6.35 million legitimate vs. ~8k fraud.



## 2. Transaction Types

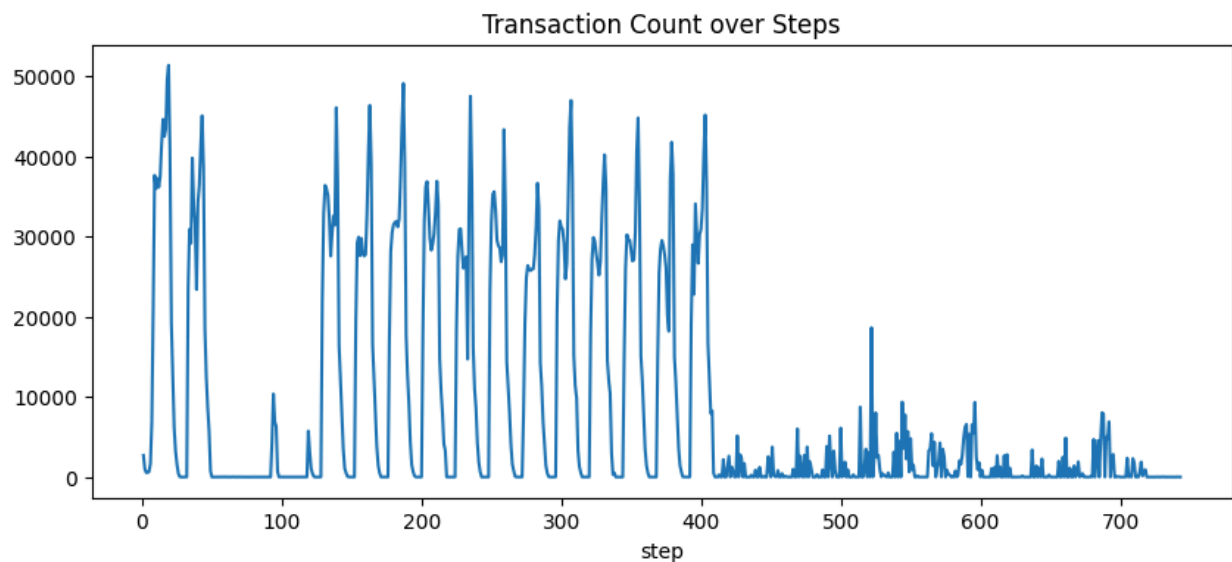
- Payment is the most common type, followed by Transfer and Cash Out.





### 3. Time Step Analysis

- Plotted transaction volume over `step`. Typically remains fairly stable but some peaks.



## 5.2 Feature Engineering

### 1. One-Hot Encode `type` :

- `type_TRANSFER` = 1 if the transaction type is "TRANSFER," else 0
- `type_CASH_OUT` , `type_DEBIT` , etc.

### 2. Delta Balances:

- `delta_balanceOrig = newbalanceOrig - oldbalanceOrig`
- `delta_balanceDest = newbalanceDest - oldbalanceDest`

### 3. Ratio:

- `amount_div_oldbalanceOrig` = `amount / oldbalanceOrig` (when `oldbalanceOrig` > 0)
- Helps detect if someone is draining their entire account in one go.

## 5.3 Model Choices and Why

We tested three unsupervised (or novelty) approaches:

### 1. IsolationForest

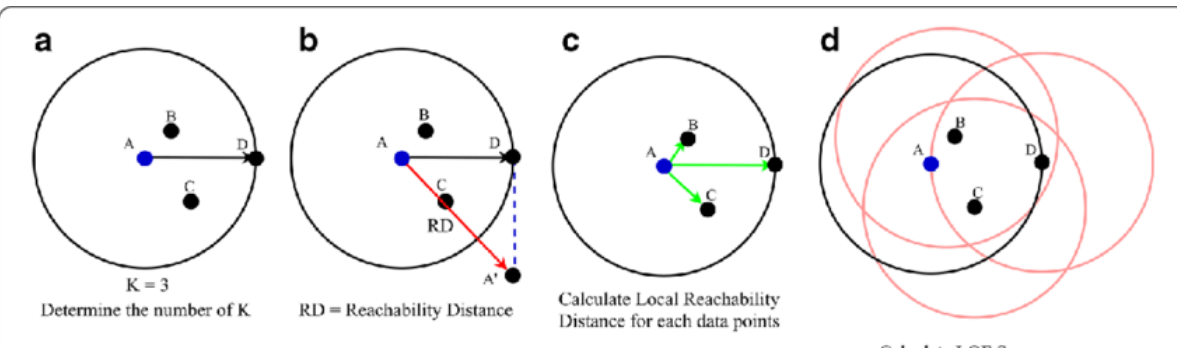

- Excels at isolating anomalies by randomly partitioning data.
- Very popular for high-dimensional data.
- In scikit-learn, we do `IsolationForest(...).fit(...)` .


### 2. Local Outlier Factor (LOF)

- Density-based approach, compares local density of each sample to neighbors.
- Good for small or moderate datasets, but can be expensive for large data (LOF is  $O(n^2)$  with naive approach).

### 3. Autoencoder (Neural Network)

- A deep learning approach.
- Trained on normal (non-fraud) data to learn a compressed representation.
- High reconstruction error on anomalies.





## 5.4 Performance of Each Model

Based on our final classification reports (since we tested them on a labeled portion), we found:

### 5.4.1 IsolationForest

- **Precision (Fraud):**  $\sim 0.01$
- **Recall (Fraud):**  $\sim 0.60$
- **F1:**  $\sim 0.02$

This means it catches  $\sim 60\%$  of fraud but with a very small precision (tons of false positives).

### 5.4.2 LOF

- **Precision (Fraud):**  $\sim 0.00$
- **Recall (Fraud):**  $\sim 0.03$  or  $0.05$

- **F1:** ~0.00 or 0.01

LOF typically had the worst performance here—lots of difficulties with large-scale data.

### 5.4.3 Autoencoder

- **Precision (Fraud):** ~0.06 or 0.07
- **Recall (Fraud):** ~0.50 or 0.54
- **F1:** ~0.11 or 0.12

This is better than the other two, giving a more balanced approach. Still plenty of false positives, but decent coverage of fraud.

--- Model Comparison ---				
	Model	Precision (Fraud)	Recall (Fraud)	F1-score (Fraud)
0	Isolation Forest	0.01	0.61	0.02
1	LOF	0.00	0.03	0.00
2	Autoencoder	0.07	0.54	0.12

## 5.5 Selecting the Final Model

We eventually picked **IsolationForest** in the real-time pipeline because:

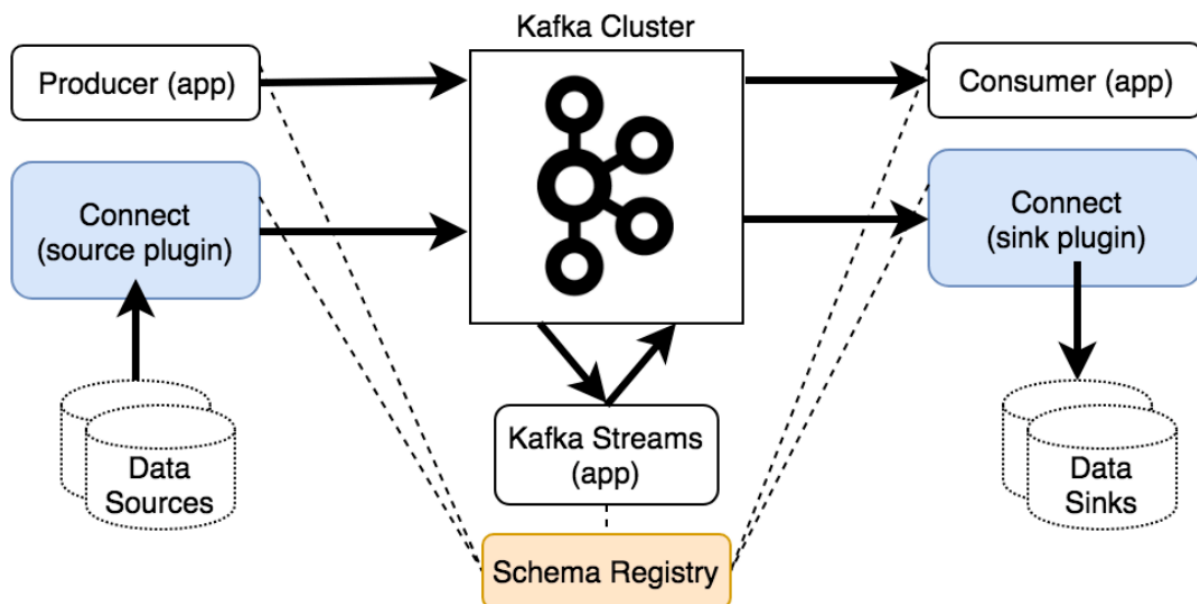
- It's relatively quick to predict.
- Doesn't require a GPU or large overhead.
- Recall was relatively strong (~60%).

Though the autoencoder had decent performance, it can be more complicated to integrate or scale. Some people might prefer the autoencoder if they can handle the overhead or want a better F1.

## 6. BUILDING THE REAL-TIME PIPELINE (KAFKA + SPARK)

### 6.1 Pipeline Architecture

1. **Kafka Producer:** A Python script ( `kafka_producer.py` ) that reads the PaySim CSV row by row, sending each transaction as a JSON message to a **Kafka topic**.
2. **Kafka Topic:** e.g. "paysim\_transactions."
3. **Spark Streaming:** A script ( `spark_streaming.py` ) that subscribes to that Kafka topic, loads the IsolationForest model ( `iso_forest.pkl` ), and runs a **UDF** to predict anomalies.
4. **Flagged Anomalies:** Spark writes them as CSV to a specified output folder.



### 6.2 Why Kafka?

- **Kafka** is a distributed messaging system, widely used for real-time streaming ingestion.
- Simulates a "live feed" of transactions.

- Allows us to scale if we had multiple producers or more consumers.

## 6.3 Why Spark Structured Streaming?

- It's a popular big-data engine for **stream** processing.
- Can do "micro-batches" in near-real time.
- Integration with Kafka is first-class ( `spark-sql-kafka-0-10` connector).

## 6.4 The `kafka_producer.py` Script

1. **Uses** `csv.DictReader` to read each row from "Transactions.csv."
2. **Transforms** row fields (like `step`, `amount`, etc.) into a JSON message.
3. **Sends** to the "paysim\_transactions" topic with `producer.send(...)`.
4. **Sleeps** ~0.01 seconds after each row to simulate ~100 transactions/sec.

## 6.5 The `spark_streaming.py` Script

### 6.5.1 Steps

1. **Create SparkSession** with Kafka support:

```
spark = SparkSession.builder \
    .appName("PaySimFraudDetection") \
    .master("local[*]") \
    .getOrCreate()
```

2. **Load Model:**

```
iso_forest = joblib.load("/home/.../iso_forest.pkl")
```

3. **Define Schema** for incoming JSON (step, amount, oldbalanceOrg, etc.).
4. **Read from Kafka:**

```
df_kafka = spark.readStream \
    .format("kafka") \
    .option("kafka.bootstrap.servers", "localhost:9092") \
    .option("subscribe", "paysim_transactions") \
    .load()
```

5. **Parse** the JSON into columns.
6. **UDF** for `predict_iso(...)` : uses `iso_forest.predict(features)` .
7. **Filter** or mark anomalies: `IF_Anomaly = 1 if pred == -1 else 0` .
8. **Write** flagged rows to CSV via:

```
df_fraud.writeStream \
    .outputMode("append") \
    .format("csv") \
    .option("path", "/home/monishaapatro/flagged_anomalies") \
    .option("checkpointLocation", "/home/monishaapatro/checkpoints") \
    .start() \
    .awaitTermination()
```

## 6.6 Google Cloud Platform

We used a **Compute Engine** VM for:

- **Installing** Kafka, Spark.
- **Running** the producer script, the Spark streaming job.
- Because it's often cheaper or straightforward than other clouds, but the same approach can be done on AWS, Azure, or local servers.

## 7. INSTALLING DEPENDENCIES AND SETTING UP



## 7.1 Dependencies

- **Java** (needed by Kafka & Spark)
- **Python 3.11**
- **Pip packages:** `pandas` , `numpy` , `scikit-learn` , `streamlit` , `joblib` , etc.
- **Spark:** downloaded from apache.org (like spark-3.5.4-bin-hadoop3).
- **Kafka:** downloaded from apache.org or used `kafka_2.13-3.5.1.tgz` .

### 7.1.1 Installing Kafka

1. Download the `.tgz` .
2. Extract to `~/kafka_folder` .
3. Start Zookeeper, then start Kafka server:

```
bin/zookeeper-server-start.sh config/zookeeper.properties
bin/kafka-server-start.sh config/server.properties
```

### 7.1.2 Installing Spark

1. Download the Spark binary.
2. Extract to `~/spark` .
3. Add environment variables or just call `~/spark/bin/spark-submit` .

## 7.2 Uploading Models and Dataset

- Transferred `iso_forest.pkl` , `Transactions.csv` , etc., to the VM using either `scp` or the GCP browser-based SSH's "Upload file."
  - Placed them in our home directory or specific folders.
  - Confirmed paths in scripts: e.g., `CSV_PATH = "/home/monishaapatro/Transactions.csv"` .
-

## 8. FLAGGED ANOMALIES AND HOW THEY ARE DETERMINED

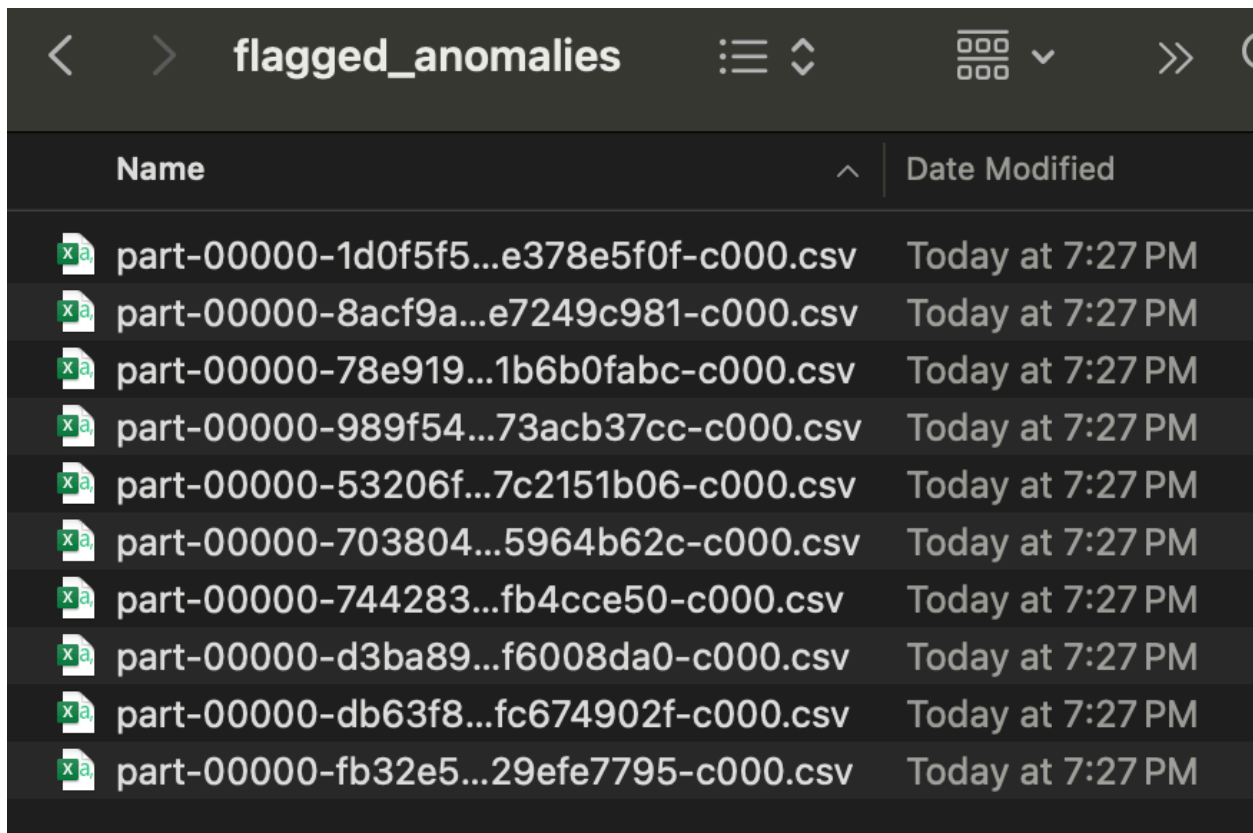
1. **IsolationForest** returns `predict(...)` = `1` for anomaly, `+1` for normal.
2. We convert `1 → 1 (IF_Anomaly=1)` meaning flagged as suspicious.
3. The script effectively says: "If anomaly, print or filter to a separate DataFrame."
4. Spark's streaming job writes those anomaly rows to **CSV** (the "flagged\_anomalies" folder).

Hence, "flagged anomalies" are simply transactions that the model deems suspicious, typically because they deviate from normal patterns.

---

## 9. TRANSPORTING ANOMALIES TO CSV

1. The final `df_anomalies` (or `df_fraud`) is written out by `df_anomalies.writeStream` with `format("csv")`.
2. Each micro-batch creates a `part-00000-xxxx.csv` file in the target folder.
3. `_spark_metadata` is also created for housekeeping.
4. This yields a growing list of CSV files each time Spark processes new data from Kafka.



Name	Date Modified
part-00000-1d0f5f5...e378e5f0f-c000.csv	Today at 7:27 PM
part-00000-8acf9a...e7249c981-c000.csv	Today at 7:27 PM
part-00000-78e919...1b6b0fab-c000.csv	Today at 7:27 PM
part-00000-989f54...73acb37cc-c000.csv	Today at 7:27 PM
part-00000-53206f...7c2151b06-c000.csv	Today at 7:27 PM
part-00000-703804...5964b62c-c000.csv	Today at 7:27 PM
part-00000-744283...fb4cce50-c000.csv	Today at 7:27 PM
part-00000-d3ba89...f6008da0-c000.csv	Today at 7:27 PM
part-00000-db63f8...fc674902f-c000.csv	Today at 7:27 PM
part-00000-fb32e5...29efe7795-c000.csv	Today at 7:27 PM

## 10. USING STREAMLIT TO PRESENT THE DATA

### 10.1 Approach

We created a **Streamlit** app ( [dashboard.py](#) ) that:

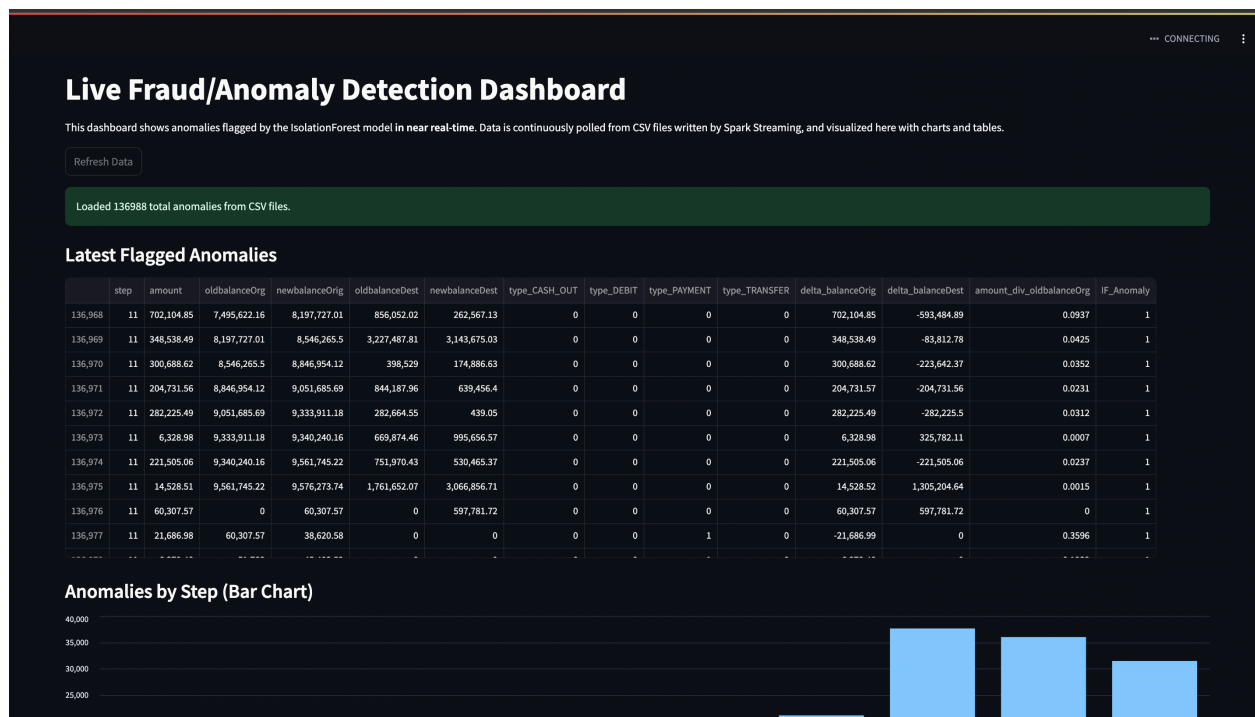
1. **Loads** the CSV files from "flagged\_anomalies/" (or whichever directory Spark is writing to).
2. **Concatenates** them to gather all flagged transactions.
3. **Visualizes** them with line charts, bar charts, tables, etc.


### 10.2 Real-Time Refresh

Because Spark is continuously writing new part-\*.csv files, the Streamlit script can periodically re-read the folder. For instance:

```
def load_anomalies_data():
    files = glob.glob("/home/monishaapatro/flagged_anomalies/*.csv")
    ...
    # read + concat data
    return df
```

A “Refresh Data” button in Streamlit might call `load_anomalies_data()` again. The user sees updated anomalies in near real time.





## 10.3 External Access

- If hosting on the **same** GCP VM, we can open the firewall for port 8501 and share the IP.
- We **deploy** to **Streamlit Community Cloud** with static sample data.

# 11. BENEFITS OF THE PROJECT

## 1. Rapid Detection:

- Instead of batch detection after the fact, suspicious transactions can be flagged within seconds.

## 2. **Scalability:**

- Kafka + Spark can handle millions of events per day.
- The unsupervised ML approach can be extended to more advanced or ensemble methods.

## 3. **End-to-End Pipeline:**

- Showcases ingestion, stream processing, ML inference, and final visualization—a full data engineering + data science solution.

## 4. **Demonstrates** skill in:

- Big data tools (Kafka, Spark)
  - Cloud environment (GCP)
  - Machine learning (IsolationForest, LOF, Autoencoder)
  - Dashboarding (Streamlit)
- 

# 12. FULL LIST OF SKILLS, TOOLS, FRAMEWORKS

- **Python** (3.11)
- **Kafka** (Apache Kafka 2.13+ distribution)
- **Spark** (Spark 3.5.4 with Structured Streaming + Kafka connector)
- **Scikit-learn** (for isolation forests, LOF)
- **TensorFlow/Keras**
- **pandas, numpy, seaborn, matplotlib** (EDA, data wrangling)
- **Streamlit** (for real-time dashboard)
- **Google Cloud Platform**
  - **Compute Engine** VM
- **Git, GitHub** (version control)
- Potentially **tmux** or **systemd** for running scripts on the VM continuously

- If using the Kaggle API, **kaggle** CLI.
- 

## 13. CONCLUSION

This project demonstrates a **complete pipeline** for **Real-Time Fraud Detection**:

1. We **ingest** synthetic transactions from the PaySim dataset in a streaming manner using Kafka.
2. **Spark** processes each transaction, applies an **IsolationForest** anomaly detection model, and flags suspicious ones.
3. These flagged anomalies are **written** to CSV.
4. A **Streamlit dashboard** reads those CSVs, producing interactive charts and tables to help investigators or analysts see suspicious transactions in near real time.
5. Everything is deployed on a GCP VM, showcasing **cloud** and **big data** skills.

Despite being synthetic, the **PaySim** dataset is large enough to **replicate** real production-scale challenges. This blueprint can adapt to a real financial system by hooking into actual transaction streams, training a more robust model on real data, and employing advanced security measures.

---

## 14. FUTURE WORK

1. **Ensemble Models**: Combine multiple anomaly detection models (e.g., autoencoder + isolation forest) for improved precision.
2. **Incorporate Additional Features**: E.g., geographic location, device info, historical user patterns.
3. **Real Production**: Use a real streaming pipeline with secure authentication, encryption, and robust monitoring.
4. **Active Learning**: Collect feedback from fraud investigators to refine the model.

5. **Advanced Visualizations:** Implement a more dynamic real-time chart that auto-refreshes every few seconds.
  6. **Auto-scaling:** Deploy Kafka + Spark on a cluster.
  7. **Deployment:** Containerize with Docker + use a platform like Cloud Run or Kubernetes for fully managed scaling.
- 

## 15. BONUS SECTIONS OR ADDITIONS

- **Performance Tuning:**
    - Spark micro-batch interval configuration, partitioning strategies for Kafka.
    - Model hyperparameter tuning for isolation forest.
  - **Security:**
    - SSL/TLS for Kafka, firewall rules for the GCP VM.
  - **Data Governance:**
    - Real deployment must respect KYC (Know Your Customer), AML (Anti-Money Laundering) regulations.
- 

## FINAL REMARKS

Built a **multi-technology** solution that merges **data science** with **data engineering**. Even though PaySim is synthetic, we have gained experience with:

- **Large-scale** data (6 million+ rows)
- **Real-time** streaming using Kafka and Spark
- **Unsupervised** ML approaches for fraud detection
- A **web dashboard** (Streamlit) for easy monitoring
- **Cloud hosting** on GCP

This sets the stage for to handle real production pipelines, adapt to more advanced fraud detection, and integrate with sophisticated data or security frameworks.